


Santa Clara University
COEN 241: Cloud Computing
HW3: Mininet
Ankita Kumar (W1620637)

The binary tree we are creating has **7 switches** and **8 hosts**:



```
root@6d75fcf106f4:~# mn --custom binary_tree.py --controller remote --topo binary_tree
*** Error setting resource limits. Mininet's performance may be affected.
*** Creating network
*** Adding controller
Unable to contact the remote controller at 127.0.0.1:6653
Connecting to remote controller at 127.0.0.1:6633
*** Adding hosts:
h1 h2 h3 h4 h5 h6 h7 h8
*** Adding switches:
s1 s2 s3 s4 s5 s6 s7
*** Adding links:
(h1, s3) (h2, s3) (h3, s4) (h4, s4) (h5, s6) (h6, s6) (h7, s7) (h8, s7) (s2, s1) (s3, s2) (s4, s2) (s5, s1) (s6, s5) (s7, s5)
*** Configuring hosts
h1 h2 h3 h4 h5 h6 h7 h8
*** Starting controller
c0
*** Starting 7 switches
s1 s2 s3 s4 s5 s6 s7 ...
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3 h4 h5 h6 h7 h8
h2 -> h1 h3 h4 h5 h6 h7 h8
h3 -> h1 h2 h4 h5 h6 h7 h8
h4 -> h1 h2 h3 h5 h6 h7 h8
h5 -> h1 h2 h3 h4 h6 h7 h8
h6 -> h1 h2 h3 h4 h5 h7 h8
h7 -> h1 h2 h3 h4 h5 h6 h8
h8 -> h1 h2 h3 h4 h5 h6 h7
*** Results: 0% dropped (56/56 received)
```

Task 1: Defining custom topologies

1. What is the output of “nodes” and “net”

```
mininet> nodes
available nodes are:
c0 h1 h2 h3 h4 h5 h6 h7 h8 s1 s2 s3 s4 s5 s6 s7
```

```
mininet> net
h1 h1-eth0:s3-eth1
h2 h2-eth0:s3-eth2
h3 h3-eth0:s4-eth1
h4 h4-eth0:s4-eth2
h5 h5-eth0:s6-eth1
h6 h6-eth0:s6-eth2
h7 h7-eth0:s7-eth1
h8 h8-eth0:s7-eth2
s1 lo: s1-eth1:s2-eth3 s1-eth2:s5-eth3
s2 lo: s2-eth1:s3-eth3 s2-eth2:s4-eth3 s2-eth3:s1-eth1
s3 lo: s3-eth1:h1-eth0 s3-eth2:h2-eth0 s3-eth3:s2-eth1
s4 lo: s4-eth1:h3-eth0 s4-eth2:h4-eth0 s4-eth3:s2-eth2
s5 lo: s5-eth1:s6-eth3 s5-eth2:s7-eth3 s5-eth3:s1-eth2
s6 lo: s6-eth1:h5-eth0 s6-eth2:h6-eth0 s6-eth3:s5-eth1
s7 lo: s7-eth1:h7-eth0 s7-eth2:h8-eth0 s7-eth3:s5-eth2
c0
```

2. What is the output of “h7 ifconfig”?

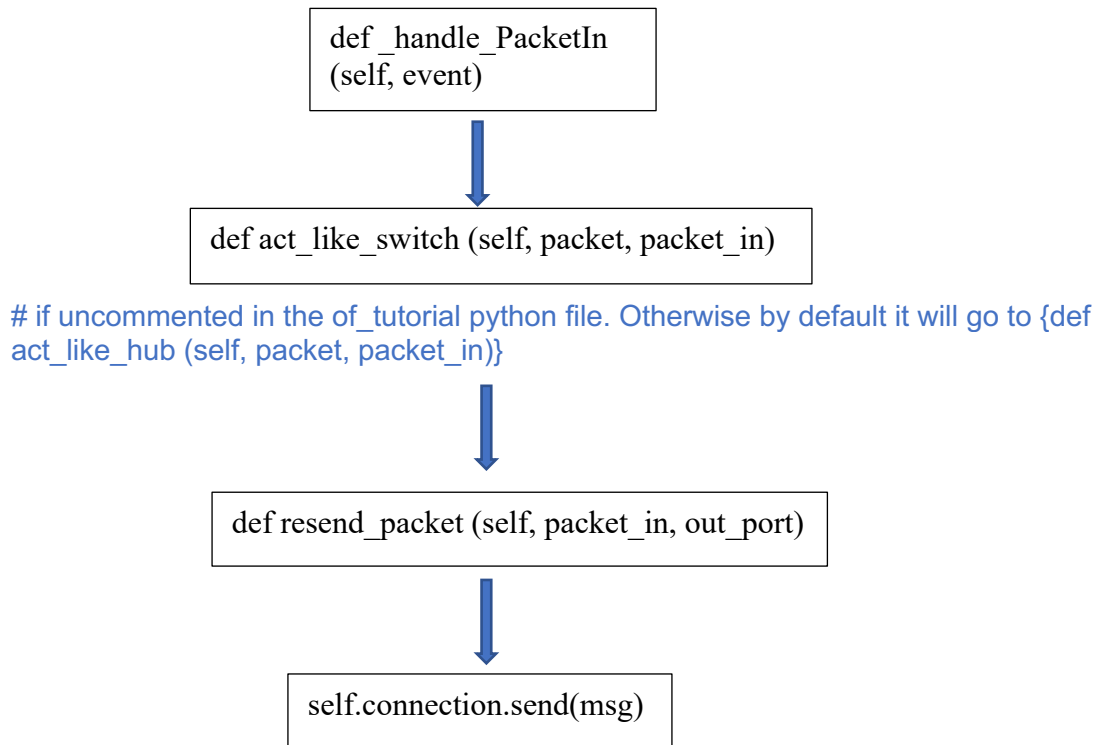
```
mininet> h7 ifconfig
h7-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.7 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::f0eb:43ff:fec8:3021 prefixlen 64 scopeid 0x20<link>
    ether f2:eb:43:c8:30:21 txqueuelen 1000 (Ethernet)
    RX packets 76 bytes 5792 (5.7 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 12 bytes 936 (936.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

Task 2: Analyze the “of_tutorial” controller

1. Draw the function call graph of this controller. For example, once a packet comes to the controller, which function is the first to be called, which one is the second, and so forth?

When packet arrives to the controller:



2. Have h1 ping h2, and h1 ping h8 for 100 times (e.g., h1 ping -c100 p2).

H1 ping h2 100 times.

```
mininet> h1 ping -c100 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=15.2 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=4.65 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=5.81 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=4.28 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=5.47 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=5.55 ms
64 bytes from 10.0.0.2: icmp_seq=7 ttl=64 time=5.92 ms
64 bytes from 10.0.0.2: icmp_seq=8 ttl=64 time=4.41 ms
```

H1 ping h8 100 times.

```
mininet> h1 ping -c100 h8
PING 10.0.0.8 (10.0.0.8) 56(84) bytes of data.
64 bytes from 10.0.0.8: icmp_seq=1 ttl=64 time=23.7 ms
64 bytes from 10.0.0.8: icmp_seq=2 ttl=64 time=18.7 ms
64 bytes from 10.0.0.8: icmp_seq=3 ttl=64 time=23.0 ms
64 bytes from 10.0.0.8: icmp_seq=4 ttl=64 time=26.6 ms
64 bytes from 10.0.0.8: icmp_seq=5 ttl=64 time=25.2 ms
64 bytes from 10.0.0.8: icmp_seq=6 ttl=64 time=27.8 ms
```

a. How long does it take (on average) to ping for each case?

- H1 ping h2 100 times:
Average time to ping: 5.527ms

```
-- 10.0.0.2 ping statistics --  
100 packets transmitted, 100 received, 0% packet loss, time 99154ms  
rtt min/avg/max/mdev = 1.644/5.527/15.235/1.337 ms
```

- H1 ping h8 100 times.
Average time to ping: 22.920 ms

```
--- 10.0.0.8 ping statistics ---  
100 packets transmitted, 100 received, 0% packet loss, time 99147ms  
rtt min/avg/max/mdev = 5.333/22.920/42.299/3.836 ms  
mininet>
```

b. What is the minimum and maximum ping you have observed?

H1 ping h2 100 times:

minimum and maximum ping observed: 1.644 ms & 15.235 ms

H1 ping h8 100 times.

minimum and maximum ping observed: 5.333 ms & 42.299 ms

c. What is the difference, and why?

H1 ping h8 took longer compared to h1 ping h2. This is due to greater number of switches between h1 & h8.

3. Run “iperf h1 h2” and “iperf h1 h8”

a. What is “iperf” used for?

It is used to test TCP bandwidth to evaluate the network performance and quality of a network line.

b. What is the throughput for each case?

“iperf h1 h2”

```
mininet> iperf h1 h2  
*** Iperf: testing TCP bandwidth between h1 and h2  
*** Results: ['9.98 Mbits/sec', '11.4 Mbits/sec']  
mininet>
```

“iperf h1 h8”

```
mininet> iperf h1 h8  
*** Iperf: testing TCP bandwidth between h1 and h8  
*** Results: ['3.09 Mbits/sec', '3.75 Mbits/sec']  
mininet>
```

- c. What is the difference, and explain the reasons for the difference?

The throughput for h1 -> h8 is low compared to h1->h2. This is because there are multiple switches between h1 & h8, as compared to h1& h2.
Hence there is more packet drop in case of h1 -> h8 as the packet flows through a greater number of switches.

4. Which of the switches observe traffic? Please describe your way for observing such traffic on switches (e.g., adding some functions in the "of_tutorial" controller).

I have added below code in the of_tutorial.py file to display traffic in the connection.

```
print("Switch:", event.connection, "source address :", packet.src, "destination  
address:", packet.dst)
```

As an output we can see there is traffic in each switch. This is in accordance to the binary tree topology created. For "iperf h1 h2", s3 observes traffic & for "iperf h1 h8", s3, s2, s1, s5, s7 observe traffic.

Task 3: MAC Learning Controller

1. Describe how the above code works, such as how the "MAC to Port" map is established. You could use a 'ping' example to describe the establishment process (e.g., h1 ping h2).

```
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=3.09 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=3.48 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=3.63 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=4.05 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=3.83 ms
^C
--- 10.0.0.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4007ms
rtt min/avg/max/mdev = 3.092/3.619/4.054/0.334 ms
mininet>
```

As "MAC to Port" map is established.

- When we enable the "act_like_switch" function, there is learning by mapping of each switch connection to a host.
- When the source MAC isn't in the map mac_to_port, script adds add the {MAC, Port} into the map mac_to_port.
- When the destination MAC is in the map, script sends the packet out the port associated with the MAC.
- When the destination MAC isn't in the map, script sends output to all ports except the input port.

2. (Comment out all prints before doing this experiment) Have h1 ping h2, and h1 ping h8 for 100 times (e.g., h1 ping -c100 p2).

h1 ping -c100 h2

```
mininet> h1 ping -c100 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=3.63 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=3.93 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=4.01 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=4.55 ms
```

h1 ping -c100 h8

```
mininet> h1 ping -c100 h8
PING 10.0.0.8 (10.0.0.8) 56(84) bytes of data.
64 bytes from 10.0.0.8: icmp_seq=1 ttl=64 time=46.6 ms
64 bytes from 10.0.0.8: icmp_seq=2 ttl=64 time=17.5 ms
64 bytes from 10.0.0.8: icmp_seq=3 ttl=64 time=18.2 ms
64 bytes from 10.0.0.8: icmp_seq=4 ttl=64 time=21.9 ms
64 bytes from 10.0.0.8: icmp_seq=5 ttl=64 time=15.7 ms
64 bytes from 10.0.0.8: icmp_seq=6 ttl=64 time=18.0 ms
```

- a. How long did it take (on average) to ping for each case?

H1 ping h2 100 times:

Average time to ping: 4.946ms

```
--- 10.0.0.2 ping statistics ---
100 packets transmitted, 100 received, 0% packet loss, time 99141ms
rtt min/avg/max/mdev = 1.179/4.946/7.733/0.880 ms
```

H1 ping h8 100 times:

Average time to ping: 18.015 ms

```
--- 10.0.0.8 ping statistics ---
100 packets transmitted, 100 received, 0% packet loss, time 99152ms
rtt min/avg/max/mdev = 3.944/18.015/46.686/4.211 ms
mininet> □
```

- b. What is the minimum and maximum ping you have observed?

H1 ping h2 100 times:

minimum and maximum ping observed: 1.179 ms & 7.733 ms

H1 ping h8 100 times:

minimum and maximum ping observed: 3.944 ms & 46.686 ms

c. Any difference from Task 2 and why do you think there is a change if there is?

Yes, there is a difference in all the three parameters:

- average ping time: Task3 < Task 2.
- There is decrease in minimum ping in task 3 compared to Task 2.
- Overall, there is a decrease in ping time compared to task2. This because of “mac_to_port “is established. There is no flooding of packets everywhere. Now the specific port to transfer the packet is learnt.

	Task 2	Task 3
<u>H1 ping h2 100 times:</u> Average time to ping	5.527ms	4.946ms
<u>H1 ping h8 100 times:</u> Average time to ping	22.920 ms	18.015 ms
<u>H1 ping h2 100 times:</u> minimum and maximum	1.644 ms & 15.235 ms	1.179 ms & 7.733 ms
<u>H1 ping h8 100 times:</u> minimum and maximum	5.333 ms & 42.299 ms	3.944ms & 46.686 ms

3. Run “iperf h1 h2” and “iperf h1 h8”.

a. What is the throughput for each case?

“iperf h1 h2”

```
mininet> iperf h1 h2
*** Iperf: testing TCP bandwidth between h1 and h2
*** Results: ['34.7 Mbits/sec', '36.5 Mbits/sec']
```

“iperf h1 h8”

```
mininet> iperf h1 h8
*** Iperf: testing TCP bandwidth between h1 and h8
*** Results: ['2.81 Mbits/sec', '3.34 Mbits/sec']
mininet> 
```

- b. What is the difference from Task 2 and why do you think there is a change if there is?

Yes, there is visible difference in the throughput for “iperf h1 h2”. It has increased in case of Task3. This is because now after the “mac_to_port” is established there is no flooding of packets everywhere. Now the specific port to transfer the packet is learnt.

Install pox:

```
root@6d75fcf106f4: ~  
root@6d75fcf106f4: ~/pox  
running python post-rtupdate hooks for python3.6...  
Processing triggers for libc-bin (2.27-3ubuntu1) ...  
root@6d75fcf106f4:~# env GIT_SSL_NO_VERIFY=true git clone https://github.com/noxrepo/pox  
Cloning into 'pox'...  
remote: Enumerating objects: 12775, done.  
remote: Total 12775 (delta 0), reused 0 (delta 0), pack-reused 12775  
Receiving objects: 100% (12775/12775), 4.84 MiB | 7.97 MiB/s, done.  
Resolving deltas: 100% (8249/8249), done.  
root@6d75fcf106f4:~# ls -lrt  
total 12  
-rw-r--r-- 1 root root 1615 Dec  3 09:00 binary_tree.py  
-rw-r--r-- 1 root root 3050 Dec  3 09:01 h1  
drwxr-xr-x 7 root root 4096 Dec  3 09:11 pox  
root@6d75fcf106f4:~# cd pox/  
root@6d75fcf106f4:~/pox# ls -lrt  
total 44  
-rwxr-xr-x 1 root root 1241 Dec  3 09:11 pox.py  
drwxr-xr-x 2 root root 4096 Dec  3 09:11 ext  
lrwxrwxrwx 1 root root 6 Dec  3 09:11 debug-pox.py -> pox.py  
-rw-r--r-- 1 root root 3069 Dec  3 09:11 README.md  
-rw-r--r-- 1 root root 892 Dec  3 09:11 NOTICE  
-rw-r--r-- 1 root root 10174 Dec  3 09:11 LICENSE  
drwxr-xr-x 16 root root 4096 Dec  3 09:11 pox  
drwxr-xr-x 2 root root 4096 Dec  3 09:11 tools  
drwxr-xr-x 3 root root 4096 Dec  3 09:11 tests  
-rw-r--r-- 1 root root 87 Dec  3 09:11 setup.cfg  
root@6d75fcf106f4:~/pox# ./pox.py log.level --DEBUG misc.of_tutorial  
POX 0.7.0 (gar) / Copyright 2011-2020 James McCauley, et al.  
DEBUG:core:POX 0.7.0 (gar) going up...  
DEBUG:core:Running on CPython (3.6.9/Jan 26 2021 15:33:00)  
DEBUG:core:Platform is Linux-5.11.0-41-generic-x86_64-with-Ubuntu-18.04-bionic  
WARNING:version:Support for Python 3 is experimental.  
INFO:core:POX 0.7.0 (gar) is up.  
DEBUG:openflow.of_01:Listening on 0.0.0.0:6633  
INFO:openflow.of_01:[00-00-00-00-00-07 2] connected  
DEBUG:misc.of_tutorial:Controlling [00-00-00-00-00-07 2]  
INFO:openflow.of_01:[00-00-00-00-00-04 3] connected  
DEBUG:misc.of_tutorial:Controlling [00-00-00-00-00-04 3]  
INFO:openflow.of_01:[00-00-00-00-00-01 4] connected  
DEBUG:misc.of_tutorial:Controlling [00-00-00-00-00-01 4]  
INFO:openflow.of_01:[00-00-00-00-00-06 5] connected  
DEBUG:misc.of_tutorial:Controlling [00-00-00-00-00-06 5]  
INFO:openflow.of_01:[00-00-00-00-00-03 6] connected  
DEBUG:misc.of_tutorial:Controlling [00-00-00-00-00-03 6]  
INFO:openflow.of_01:[00-00-00-00-00-02 7] connected  
DEBUG:misc.of_tutorial:Controlling [00-00-00-00-00-02 7]  
INFO:openflow.of_01:[00-00-00-00-00-05 8] connected  
DEBUG:misc.of_tutorial:Controlling [00-00-00-00-00-05 8]  
DEBUG:openflow.of_01:1 connection aborted
```